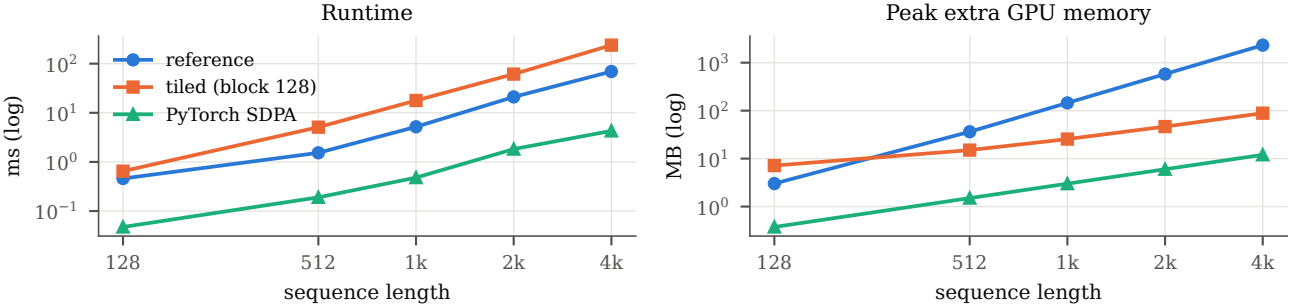
Runtime
reference
tiled (block 128)
PyTorch SDPA
ms (log)
10^2
10^1
10^0
10^{-1}
128
512
1k
2k
4k
sequence length
Peak extra GPU memory
MB (log)
10^3
10^2
10^1
10^0
128
512
1k
2k
4k
sequence length