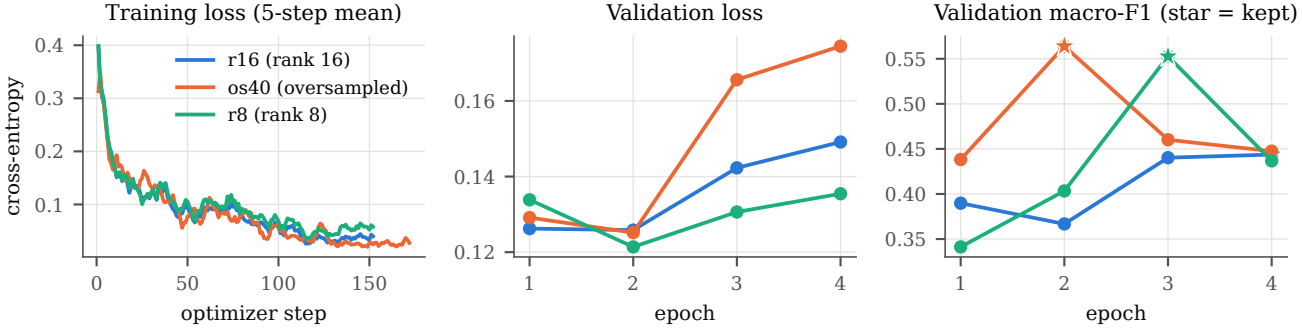
Training loss (5-step mean)
cross-entropy
0.4
0.3
0.2
0.1
r16 (rank 16)
os40 (oversampled)
r8 (rank 8)
0
50
100
150
optimizer step
Validation loss
0.16
0.14
0.12
1
2
3
4
epoch
Validation macro-F1 (star = kept)
0.55
0.50
0.45
0.40
0.35
1
2
3
4
epoch